

Delta Sharing


Learning Objectives

- ▶ What is Delta Sharing
- ▶ How it works under the hood
- ▶ Costs and limitations


Delta Sharing

- ▶ Open protocol for secure data sharing
- ▶ Share live data with organizations regardless of computing platforms


How it Works



Under the Hood



Delta Sharing on Databricks



Delta Sharing on Databricks

- ▶ D2D: The Databricks-to-Databricks sharing protocol
 - ▶ Share data between Databricks clients
 - ▶ Collection of tables, views, volumes, and notebooks
 - ▶ Leverages built-in authentication with no token exchange
- ▶ D2O: The Databricks open sharing protocol
 - ▶ Share data with users on any computing platform
 - ▶ Share tables only
 - ▶ Requires authentication via bearer tokens or OpenID Connect (OIDC) Federation

History Sharing

- ▶ Databricks-to-Databricks sharing can allow recipients to perform time travel queries and streaming reads by enabling history sharing:
 - ▶ **ALTER SHARE <share_name> ADD TABLE <table_name> WITH HISTORY;**
- ▶ Allows querying the change data feed of the shared table.
 - ▶ CDF must be enabled on the table before you share it WITH HISTORY.
- ▶ Improves performance by leveraging temporary cloud storage credentials scoped-down to the root directory of the shared table.
 - ▶ Resulting in performance equivalent to direct access to the source table.
 - ▶ Performance benefits do not apply to partitioned tables.

Roles

- ▶ Users authorized to create and manage Delta Shares:
 - ▶ Metastore admins
 - ▶ Users with CREATE SHARE privilege on the metastore

Costs

- ▶ Delta Sharing does not require data replication
- ▶ Egress cost
 - ▶ Within a region: no egress cost
 - ▶ Cross-clouds or cross-regions: cloud vendor charges data egress fees, instead:
 - ▶ Clone the shared data to local regions
 - ▶ Share data from Cloudflare R2

Limitations

- ▶ Read-Only Access Model
- ▶ Data Format Constraints
 - ▶ Only Delta tables are supported for sharing